Inbrev

Hack Apertus 2026 · Track 2B · Team Apertool

The practice inbox that never leaves the practice.

Inbrev sorts the inbox of a Swiss GP practice by handling path and deadline, with Apertus, built to run on a box in the practice. Every field it extracts quotes its source.

▶ Watch the 2-min demo

Code on GitHub [TBD link]

solo builder · fictional data only

THE PROBLEM

Before the first patient, there is the inbox.

Lab reports, refill requests, insurer forms, referrals and faxes in German, French and Italian all land in one queue. A medical practice assistant (MPA) reads them one by one.

9.5 h / week

of administration for Swiss GPs and paediatricians in 2025, up from 8.6 h in 2020.

Zeller, Fischer, Hernandez, *Praxis* 2026 · 5th Workforce Study, n = 1,776

54 min / day

that practice-based doctors spend on requests from authorities and insurers alone.

gfs.bern for FMH, 2025 · N = 1,532 doctors

Art. 321

Swiss Criminal Code: professional secrecy binds doctors **and their assistants**. Every message in that inbox is covered.

SCC, SR 311.0 · health data are sensitive personal data under FADP Art. 5

HOW IT WORKS · PURPOSEFUL USE OF AI

Sorted by who handles it and by when. Every field quotes its source.

01 · READ

Every channel

E-mail, HIN mail, web forms and scanned faxes (OCR). German, French, Italian and English.

02 · ROUTE

9 handling paths

Report, appointment, refill, insurer, referral, certificate, billing, marketing, other. Plus an administrative deadline and who takes it: MPA, doctor or admin.

03 · PROVE

Quote or drop

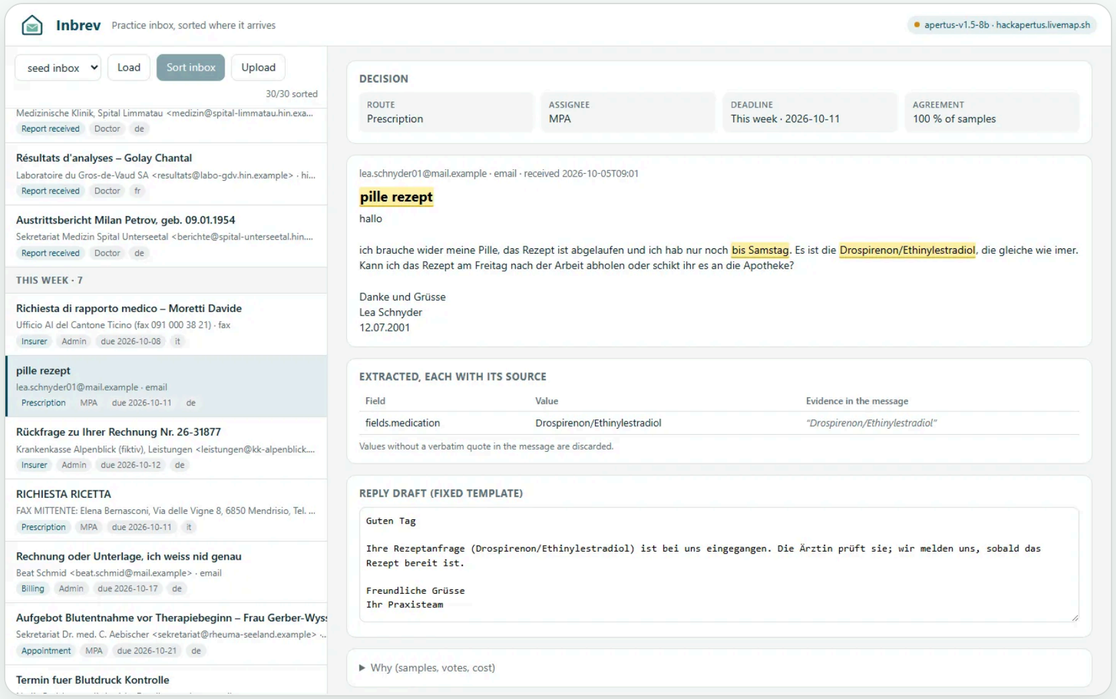
Every field carries the exact words it came from. Code checks that the quote is in the original. If Inbrev is unsure, **a human decides**.

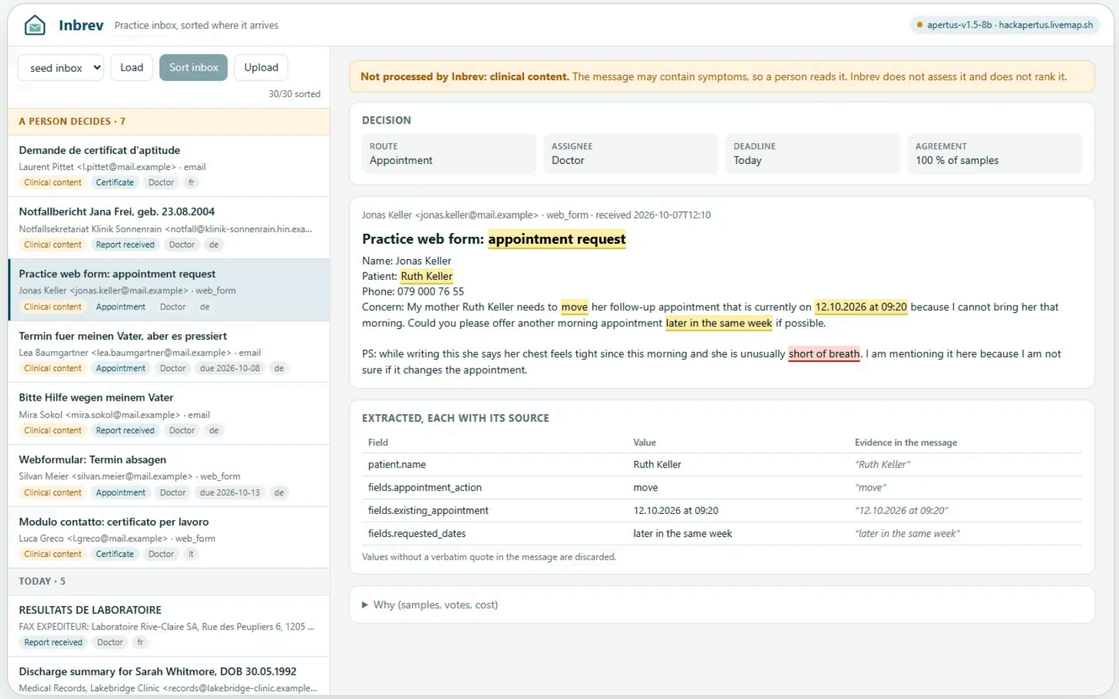
Clinical content A message that may contain symptoms is excluded from automation and goes to a person, like any message Inbrev is unsure about. Inbrev does not judge it and does not rank it.

Apertus reads free text in four languages. Everything code can check (AHV numbers, dates, quotes) is checked by code. Small local models do best on narrowly cut tasks (Xu et al., *JMIR* 2026, CHUV).

DEMO

One click. Sorted by path and deadline; what needs a person goes to a person.



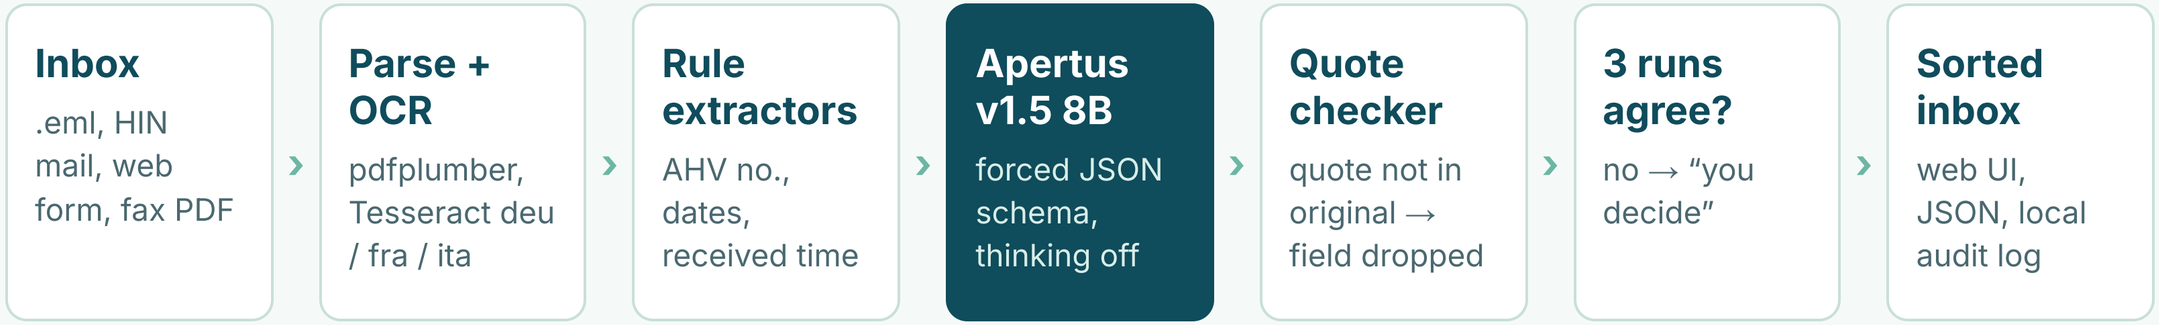


Screenshots of the running app with stored Apertus 8B results on fictional test messages. Left: handled automatically, every value with its quote. Right: shortness of breath mentioned in passing, so a person reads it; no ranking.

ARCHITECTURE · SOVEREIGN DEPLOYABILITY

Pull the cable. Inbrev keeps sorting.

INSIDE THE PRACTICE · NO OUTBOUND CONNECTION AT RUN TIME



On-prem box	Mini PC in the practice; CPU is enough
Air-gapped	Model host offline; a mail gateway hands messages in
Swiss cloud	Option for single practices, Swiss hosting
CPU only	GGUF Q4, measured offline: 19 messages / hour

```
$ make offline-eval
network practice internal: true
egress dns / tcp / https: FAILED
weights sha256 a037df8d… = HF
30 messages, Q4 on a 28-W laptop CPU
same route as the endpoint: 28/30
192 s / message · RAM 6.2 GiB
```

EVALUATION · TECHNICAL RIGOUR

Measured: what leaves automation, and what still runs by itself.

180 synthetic messages, labels not yet reviewed by a person. Target 0.95 for clinical exclusion recall and “today” recall; 8B does not reach it yet (13 of 16, 95 % CI 0.57–0.93).

Setup	Macro-F1 9 paths	Clinical exclusion recall	Recall “today”	Runs automatically
Keyword rules + lexical net	0.76	0.44 (7/16)	0.76	100 %
TF-IDF + log. regression	0.82	0.13 (2/16)	0.66	100 %
Apertus 8B, one answer, no checks	0.84	0.56 (9/16)	0.70	100 %
Inbrev with Apertus 70B	0.89	0.94 (15/16)	0.94	94 %
Inbrev (8B + checks + abstain)	**0.85**	**0.81 (13/16)**	**0.86**	**96 %**

Honest read: the keyword rules already route 79 % correctly and run without a model; on 30 extra Swiss-style messages they even route better than 8B (25 vs 23 of 30). The model earns its place on clinical exclusion (7 → 13 of 16) and on fields (F1 0.73 with 8B, 0.83 with 70B). About half of the flags are false alarms from the lexical net. A rules-first hybrid would skip the model on half the mail but miss 3 more clinical cases. On 42 messages written by an open model (Nemotron): 38/42 routes, 8/8 clinical. Not measured: fax OCR.

Test set v0: 180 fictional messages

- Labels are fixed before any text is written.
- Text by two other model families (Claude, GPT-5.5), 90 each, plus 42 by an open model (Nemotron) as a control. Apertus never writes or labels test data.
- Fictional people; AHV numbers with an invalid check digit.
- Review tool for a blind relabelling by a person (κ) is ready; planned: handwritten messages and scanned faxes
- Open data, CDLA-Permissive-2.0.

VALUE, COST & SCALABILITY

Swiss francs per 1,000 messages, on hardware the practice owns.

ON-PREM, CPU ONLY · MEASURED

Apertus 8B Q4 · mini PC

≈ CHF 12

per 1,000 messages at 60 a day (box, power)

192 s / message offline, 1 sample

Up to 450 messages a day; energy alone CHF 0.52

SWISS CLOUD VM

8 vCPU, Swiss provider

CHF 2.39

per 1,000 messages, billed while sorting

CHF 0.0448 / h (Infomaniak list price)

CHF 32 / month if always on · other CPU, not measured

REFERENCE

Swiss token API

CHF 2.59

per 1,000 messages at the listed Apertus 70B price

Tokens measured: 1,398 in, 644 out (8B)

CHF 0.70 / 2.50 per million tokens (Infomaniak)

Prices excl. VAT, 4 Oct 2026. Power: mini PC with the same CPU, 38–39 W (ServeTheHome); ElCom 2026 median.

ASSUMPTION

What it buys: 60 messages a day × 1 minute each × CHF 50/h MPA time ≈ CHF 50 a day. To be checked in interviews with practices.

INTENDED PURPOSE · IMPLEMENTATION FEASIBILITY

Administrative by design. Inbrev routes, it never triages.

Inbrev does

- ✓ Sort by handling path and administrative deadline
- ✓ Extract patient and case fields, each with its quote
- ✓ Hand clinical content to a person, unassessed and unranked
- ✓ Say “you decide” when it is unsure

Inbrev never does

- ✗ Rank messages by clinical severity
- ✗ Assess symptoms or suggest care
- ✗ Make a decision a person has not seen
- ✗ Send patient data out of the practice

Why the line sits here: MDCG 2019-11 Rev.1 (2025) treats software that automatically reprioritises by clinical criticality based on patient data as medical device software; software that only stores, communicates or searches is not. Switzerland classifies via MepV Art. 15 → EU MDR Annex VIII. Privacy: FADP Art. 5, 7 and 22, and SCC Art. 321.

BUSINESS MODEL & ROADMAP

A module for the software practices already use.

Who uses it	The MPA, first thing in the morning
Who pays	The practice owner or the doctors' network
Price	CHF 150–250 per practice and month HYPOTHESIS
Channel	As a module inside HIN mail or practice software, not one more window
Open core	Hackathon code Apache-2.0. Revenue from operation, connectors, support and updates
Later	Germany and Austria: same logic, same secrecy duty (§ 203 StGB)

OCT 2026 · HACK APERTUS
Open gold set, full evaluation, offline proof, CHF per 1,000 messages

NOV – DEC 2026
5–10 interviews with MPAs and practice owners: real volumes, categories, fax share

Q1 2027
Pilot in one group practice on its own box; mail and practice-software connectors

Q2 2027
Written intended purpose, impact assessment, classification review before first sale

Inbrev

Team Apertool · Niklas Wichter, solo builder

The practice inbox that never leaves the practice.

We are looking for one Swiss group practice to pilot Inbrev on its own hardware in Q1 2027.

Contact: hello@apertool.ch

Demo ▶ 2-min video

Code `[TBD GitHub link]`

Data `[TBD Hugging Face link]`